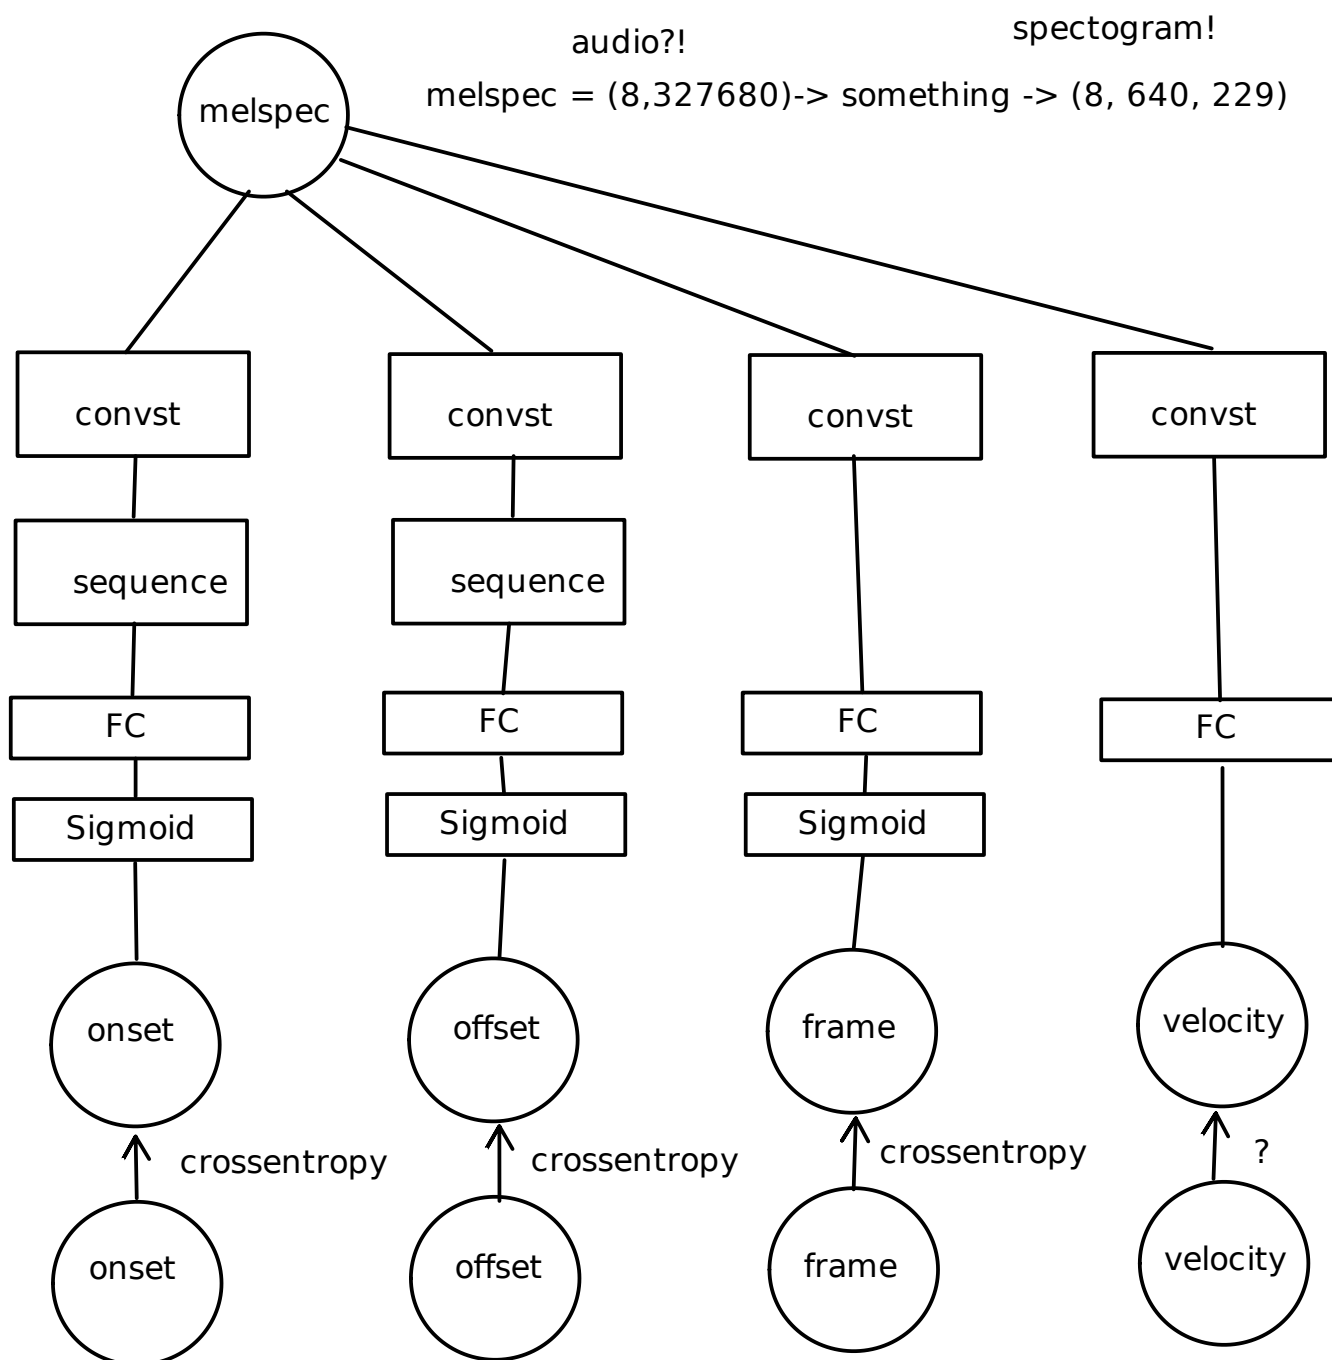


Input: mel (8,640,229)

4 Predictions: (8,640,88) onset, offset, frame, velocity ; same shapes as labels
(batch,frames,input_features)



Sequence = nn.LSTM(input_features, recurrent_features, batch_first=True, bidirectional=True)

altered for prediction to support longer sequences!

where are the hidden states?